


# Introduction to OpenMP

# Introduction

- OpenMP is designed for **shared memory** systems.
- OpenMP is easy to use
  - achieve parallelism through compiler directives
  - or the occasional function call
- OpenMP is a “quick and dirty” way of parallelizing a program.
- OpenMP is usually used on existing serial programs to achieve moderate parallelism with relatively little effort

# Computational Threads

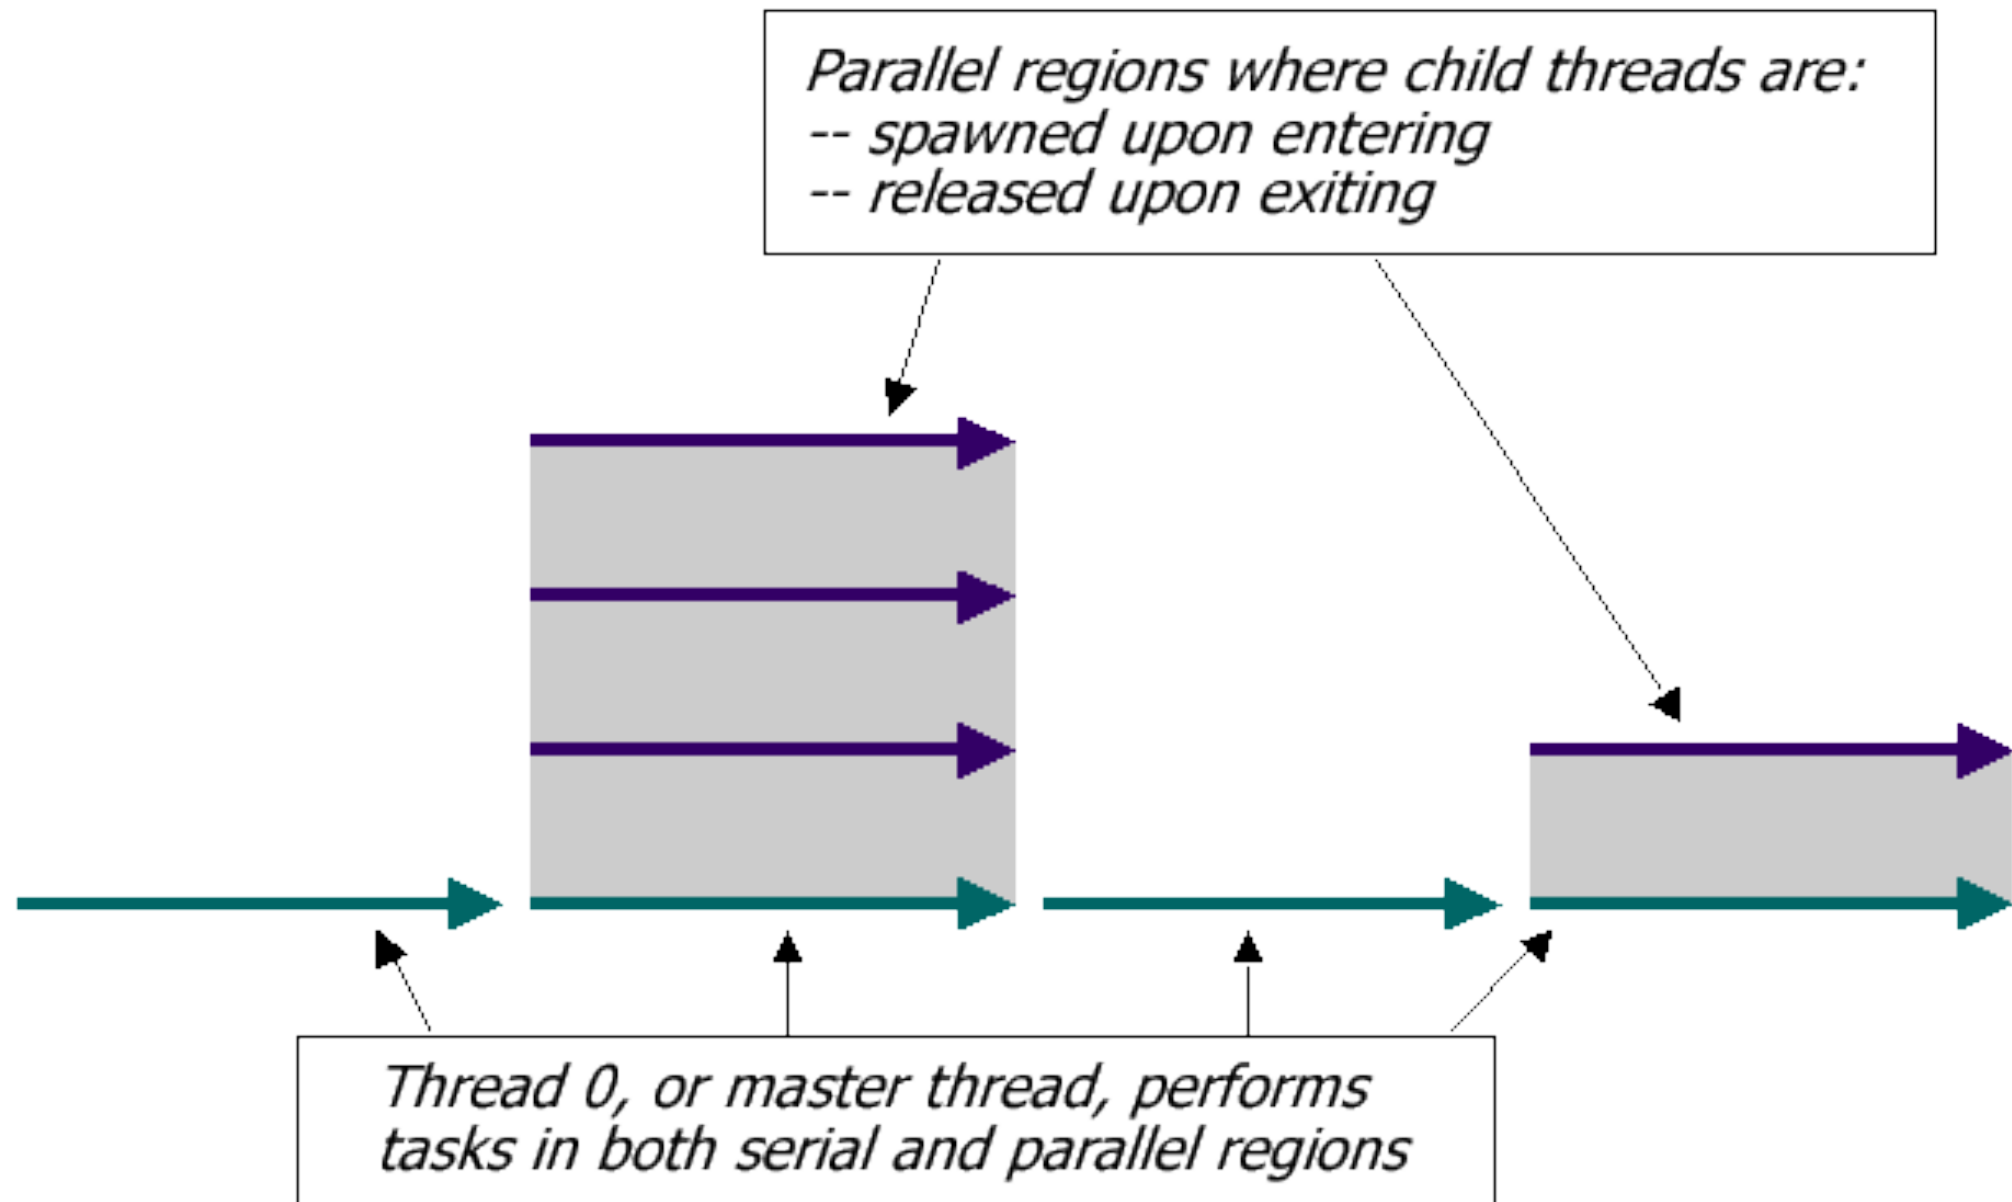
- Each processor has one thread assigned to it
- Each thread runs one copy of your program



# OpenMP Execution Model

- In **OpenMP**, execution begins only on the master thread. Child threads are spawned and released as needed.
  - Threads are spawned when program enters a **parallel region**.
  - Threads are released when program exits a **parallel region**

# OpenMP Execution Model



# Parallel Region Example: For loop

## Fortran:

```
!$omp parallel do  
do i = 1, n  
  a(i) = b(i) + c(i)  
enddo
```

This comment or pragma tells openmp compiler to spawn threads \*and\* distribute work among those threads

These actions are combined here but they can be specified separately between the threads

## C/C++:

```
#pragma omp parallel for  
for(i=1; i<=n; i++)  
  a[i] = b[i] + c[i];
```

# Pros of OpenMP

- Because it takes advantage of shared memory, the programmer does not need to worry (that much) about data placement
- Programming model is “serial-like” and thus conceptually simpler than message passing
- Compiler directives are generally simple and easy to use
- Legacy serial code does not need to be rewritten


# Cons of OpenMP

- Codes can only be run in shared memory environments!
  - In general, shared memory machines beyond ~8 CPUs are much more expensive than distributed memory ones, so finding a shared memory system to run on may be difficult
- Compiler must support OpenMP
  - whereas MPI can be installed anywhere
  - However, gcc 4.2 now supports OpenMP



# Cons of OpenMP

- In general, only moderate speedups can be achieved.
  - Because OpenMP codes tend to have serial-only portions, Amdahl's Law prohibits substantial speedups
- Amdahl's Law:
  - $F$  = Fraction of serial execution time that cannot be parallelized
  - $N$  = Number of processors



Execution time = 
$$\frac{1}{F + (1 - F)/N}$$

If you have big loops that dominate execution time, these are ideal targets for OpenMP

# Compiling and Running OpenMP

- True64: **-mp**
- SGI IRIX: **-mp**
- IBM AIX: **-qsmp=omp**
- Portland Group: **-mp**
- Intel: **-openmp**
- gcc (4.2) **-fopenmp**

# Compiling and Running OpenMP

- OMP\_NUM\_THREADS environment variable sets the number of processors the OpenMP program will have at its disposal.
- Example script

```
#!/bin/tcsh
```

```
setenv OMP_NUM_THREADS 4
```

```
mycode < my.in > my.out
```

# Sections: Functional parallelism

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        block1
        #pragma omp section
        block2
    }
}
```




Image from: <https://computing.llnl.gov/tutorials/openMP>

# Parallel DO/for: Loop level parallelism

Fortran:

```
!$omp parallel do  
do i = 1, n  
  a(i) = b(i) + c(i)  
enddo
```

C/C++:

```
#pragma omp parallel for  
for(i=1; i<=n; i++)  
  a[i] = b[i] + c[i];
```




Image from: <https://computing.llnl.gov/tutorials/openMP>

# OpenMP Functions

- **omp\_set\_num\_threads(n)**
  - sets number of openmp threads to n
- **omp\_get\_num\_threads()**
  - returns integer value of total number of threads
- **omp\_get\_thread\_num()**
  - returns individual thread id number.

# OpenMP *single* clause

```
#include<stdio.h>
#include<omp.h>

int main()
{
    int num_thds, myid;

    omp_set_num_threads(4);

    #pragma omp parallel
    {
        num_thds = omp_get_num_threads();

        myid = omp_get_thread_num();

        #pragma omp single
        printf("\nHello World from thd num %d out of %d thds!", myid,
num_thds);
    }
    printf("\nProgram Exit!\n");
}
```

# OpenMP *master* clause

```
int main()
{
    int num_thds, myid;

    //omp_set_num_threads(4);

    #pragma omp parallel private(num_thds, myid)
    {
        num_thds = omp_get_num_threads();

        myid = omp_get_thread_num();

        #pragma omp master
        {
            printf("\nI am Master: %d out of %d thds!", myid, num_thds);
        }

        printf("\nAll: %d out of %d thds!", myid, num_thds);
    }
    printf("\nProgram Exit!\n");
}
```



# OpenMP *private* clause

```
int main()
{
    int num_thds, myid;
    int data = 10;

    omp_set_num_threads(4);

    #pragma omp parallel private(num_thds, myid, data)
    {
        num_thds = omp_get_num_threads();

        myid = omp_get_thread_num();

        data = data + myid;

        printf("\nSection 1: From thd num %d out of %d thds : data = %d", myid,
num_thds, data);
    }

    printf("\n\ndata = %d \n", data);
}
```

# OpenMP *firstprivate* clause

```
int main()
{
    int num_thds, myid;
    int data = 10;

    omp_set_num_threads(4);

    #pragma omp parallel private(num_thds, myid) firstprivate(data)
    {
        num_thds = omp_get_num_threads();

        myid = omp_get_thread_num();

        data = data + myid;

        printf("\nSection 2: From thd num %d out of %d thds : data = %d", myid,
num_thds, data);
    }

    printf("\n\n data = %d \n", data);
}
```

# OpenMP *threadprivate* clause

```
omp_set_num_threads(4);

#pragma omp threadprivate(val)

#pragma omp parallel
{
    num_thds = omp_get_num_threads();

    myid = omp_get_thread_num();

    val = 50;

    printf("\nSection 1: from thd num %d out of %d thds : val = %d", myid, num_thds, val);
}

printf("\n");

#pragma omp parallel
{
    myid = omp_get_thread_num();

    val = val + myid;
}

printf("\n");

#pragma omp parallel
{
    num_thds = omp_get_num_threads();

    myid = omp_get_thread_num();

    printf("\nSection 2: from thd num %d out of %d thds : val = %d", myid, num_thds, val);
}
```

# OpenMP ***barrier*** clause

```
#include<stdio.h>
#include<omp.h>

int main()
{
    int num_thds, myid;

    omp_set_num_threads(4);

    #pragma omp parallel
    {
        num_thds = omp_get_num_threads();

        myid = omp_get_thread_num();

        printf("\nFirst printf: %d out of %d thds!", myid, num_thds);

        #pragma omp barrier

        printf("\nSecond printf: %d out of %d thds!", myid, num_thds);

    }
    printf("\nProgram Exit!\n");
}
```

# OpenMP *shared, critical* clause

```
#include<stdio.h>
#include<omp.h>

//Gets the total of all the myid's

int main()
{
    int myid, total;

    omp_set_num_threads(4);

    #pragma omp parallel private(myid) shared(total)
    {
        myid = omp_get_thread_num();

        #pragma omp atomic update
        total += myid;
    }
    printf("\n total = %d", total);
    printf("\nProgram Exit!\n");
}
```

# OpenMP *reduction* clause

## ***Addition***

```
#pragma omp parallel for default(none) shared(A) reduction(+:sum)
    for(i=0;i<ARRSIZE;i++)
    {
        sum = sum + A[i];
    }
```

## ***Multiplication***

```
#pragma omp parallel for default(none) shared(A) reduction(*:val)
    for(i=0;i<ARRSIZE;i++)
    {
        val = val * A[i];
    }
```

# OpenMP *reduction* clause

## *Min*

```
#pragma omp parallel for default(none) shared(A) reduction(min:val)
for(i=0;i<ARRSIZE;i++)
{
    if(val > A[i])
    {
        val = A[i];
    }
}
```

## *Max*

```
#pragma omp parallel for default(none) shared(A) reduction(max:val)
for(i=0;i<ARRSIZE;i++)
{
    if(val < A[i])
    {
        val = A[i];
    }
}
```


# OpenMP *section* clause

```
omp_set_num_threads(2);
#pragma omp parallel sections
{
    #pragma omp section
    {
        num_thds = omp_get_num_threads();
        myid = omp_get_thread_num();
        printf("\nSection 1: thd num %d out of %d thds!", myid, num_thds);
    }
    #pragma omp section
    {
        num_thds = omp_get_num_threads();
        myid = omp_get_thread_num();
        printf("\nSection 2: thd num %d out of %d thds!", myid, num_thds);
    }

    #pragma omp section
    {
        num_thds = omp_get_num_threads();
        myid = omp_get_thread_num();
        printf("\nSection 3: thd num %d out of %d thds!", myid, num_thds);
    }
}
printf("\nProgram Exit!\n");
```



# Steady State Heat Conduction




## Steady State Heat Conduction

Phenomenon is modelled using Laplace's equation

$$\frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} = \nabla^2 T = 0$$

which can be discretised on a grid as,



## Steady State Heat Conduction

The discretised equation at a single point (i,j) is

$$\frac{T_{i-1,j} - T_{i,j} + T_{i+1,j}}{(\Delta x)^2} + \frac{T_{i,j-1} - T_{i,j} + T_{i,j+1}}{(\Delta y)^2} = 0$$

Assemble all the equations for all unknown points in the Matrix form and then solve

$$Ax = B$$

You can choose any Linear Algebra Solver (iterative or Direct).  
Iterative is more efficient.